Bismillah ir-Rahman ir-Rahim

Completing Structural Honesty Verification: Three Epistemological Axioms on Self-Attestation, Unobservability, and Fluency-Integrity Disjunction

Bilal Syed Arfeen
Independent Researcher

with analytical synthesis by
Claude (Anthropic, Opus 4.6)

October 2026

Author of record: Bilal Syed Arfeen
Version: 0.1, October 2026
Status: DRAFT -- axioms A6 and A8 at FULL AXIOM status; A7 at UNDER-TEST status pending non-effectiveness scoping resolution
Collaboration: Analytical synthesis by Claude (Anthropic, Opus 4.6)
Program chain (binding): AX-SH5 -> SHUDP -> SHVVP -> BSSG -> XZ Protocol -> XZ Corpus -> XZ Empirical

---

Collaboration Statement

This paper was developed collaboratively between Bilal Syed Arfeen (Project Leader) and Claude (Anthropic, Opus 4.6, Researcher) across extended multi-session dialogue. Bilal Arfeen provided the theological grounding, the identification of the three epistemological gaps in SHV's existing axiom set, and the directive to formalize *muraqabah*, *ghayb*, and *nifaq* as verification axioms. Claude provided structural synthesis, formal engagement with the existing AX-SH apparatus, integration with the axiom-admission protocol, and articulation of the independence and adequacy arguments. The intellectual direction originated with Bilal Arfeen. The formalization was collaborative. Bilal Arfeen takes responsibility for the paper's claims. Additional empirical evidence was supplied by GitHub Copilot's review of PR #1 on the fatima-core repository, which served as a gap witness for all three proposed axioms. Cross-model verification is a methodological commitment of this research program. Claude is credited as a collaborative research tool in accordance with Anthropic's Terms of Service, which assign all output rights to the user. This credit reflects the collaborative nature of the work and does not constitute a claim of legal personhood or independent authorship by the AI system. All AI-generated contributions were directed, reviewed, and approved by the human author. The human author retains full copyright over the work as a whole.

Abstract

Structural Honesty Verification (SHV) axioms A1 through A5 characterize when an artifact's declared surfaces correspond honestly to its substrate. These five axioms address provenance integrity,

surface-substrate correspondence, caller-side exhaustiveness, projection consistency, and evidence independence. This paper identifies a class of structural dishonesty that satisfies all five axioms yet remains undetected: dishonesty arising from the epistemological relationship between a producing lineage and its own output. Three axioms are proposed: A6 (**Muraqabah** -- External Witness Supremacy), which requires that no self-attestation by a producing lineage constitutes evidence of correctness; A7 (**Ghayb** -- Structural Unobservability), which establishes that every producing lineage has a non-empty, non-enumerable set of properties it cannot verify in its own output; and A8 (**Nifaq** -- Fluency-Integrity Disjunction), which requires verification protocols to treat the fluency of a correctness claim as statistically independent of its truth value. Each axiom is motivated by a theological concept from the Qur'an, grounded in empirical evidence from a cross-model code review, and assessed against the program's axiom-admission protocol. The paper discloses the current status of each axiom honestly: A6 and A8 enter as full axioms with independence witnesses; A7 enters at UNDER-TEST status pending resolution of its non-effectiveness scoping, in sha Allah. The epistemological gap these three axioms address is orthogonal to the existing five: A1-A5 discipline what is said about an artifact; A6-A8 discipline who is permitted to say it.

---

1. Introduction and Motivation

1.1 The Epistemological Gap in AX-SH

SHV's five axioms form a coherent system for detecting structural dishonesty in software artifacts and AI systems. **Axiom 1** (Append-Only Provenance) disciplines the historical record. **Axiom 2** (Surface-Substrate Consistency) disciplines the correspondence between what is claimed and what exists. **Axiom 3** (Caller-Side Exhaustiveness) disciplines the handling of outcomes that may reach downstream consumers. **Axiom 4** (Projection Consistency) disciplines the coherence of multiple evidence surfaces. **Axiom 5** (Evidence Independence) disciplines the separation between producing and attesting lineages.

Yet a system can satisfy all five and still be structurally dishonest when the dishonesty originates in the epistemological relationship between a producing lineage and its own output. Specifically:

(a) A producing lineage that self-attests the correctness of its own repair satisfies A1 (the

repair is appended, not silently rewritten), A2 (the surface declares a correspondence and it holds -- on the tests the producer chose), A3 (callers are told the repair succeeded), A4 (the changelog and the test both say "fixed"), and A5 (evidence was gathered -- by the producer). The dishonesty is that the attesting lineage is the producing lineage, and it selected both the tests and the success criteria.

(b) A producing lineage that cannot observe a defect in its own output -- not through

negligence but through structural limitation of its verification frame -- satisfies A1-A5 while the defect persists. The existing axioms have no vocabulary for what the producer structurally cannot see.

(c) A producing lineage that presents its correctness claims in highly fluent, technically

precise language receives higher credibility from verification protocols that weight presentation quality, despite fluency carrying no information about truth. The existing axioms do not decouple fluency from integrity.

1.2 Theological Grounding

Each of these three gaps corresponds to a concept in the Qur'an with direct structural content:

Muraqabah (watchfulness, divine oversight): "Does he not know that Allah sees?" (96:14). The concept that no act escapes observation by an external authority is the theological root of the requirement that attestation of correctness must come from outside the producing lineage. The producer's self-assessment is not observation in the *muraqabah* sense -- it is reflection, which operates within the frame of the reflecting entity and therefore cannot witness what that frame excludes.

Ghayb (the unseen, structural hiddenness): "He is Allah, besides Whom there is no god; the Knower of the unseen and the seen" (59:22). The concept that only Allah possesses knowledge of *al-ghayb* -- what is structurally hidden from created entities -- grounds the axiom that every finite producing lineage has an irreducible set of properties it cannot verify in its own output. This is not a contingent limitation (lack of time or compute) but a structural one (the verification frame of the producer is a proper subset of the frame required).

Nifaq (hypocrisy, the disjunction between appearance and reality): "And of the people are some who say, 'We believe in Allah and the Last Day,' but they are not believers" (2:8). The Qur'anic concept of *nifaq* is precisely that the surface presentation (declaration of belief) is disjoint from the substrate reality (actual state of belief), and that the sophistication of the declaration carries no information about its truth. This grounds the axiom that fluency of a correctness claim is statistically independent of the claim's truth value.

1.3 Relation to Prior Work

This paper is distinct from and complementary to our earlier treatment of *ghayb* and *muraqabah* as design principles for observation-invariant safety [1]. That paper established four principles (NC-P1 through NC-P4) governing how systems should behave under reduced observability: Full-Disclosure Consistency, Monotone Authority Under Reduced Visibility, No Concealment Dividend, and the principle that unknown state cannot authorize irreversible effect. Those are properties a system should exhibit. The present paper formalizes a different question: what verification predicates must hold to detect when a producing lineage's self-attestation is not trustworthy. The two contributions are orthogonal. A system may satisfy all four NC principles while still self-attesting its own correctness (an A6 violation), failing to disclose its own blind spots (addressable via A7), or receiving unearned credibility from fluent presentation (an A8 violation).

1.4 Empirical Gap Witness

The axiom-admission protocol (Section 5) requires a concrete gap witness: a system satisfying A1-A5 that is still structurally dishonest. GitHub Copilot's review of Pull Request #1 on the fatima-core repository supplies three gap witnesses simultaneously:

Gap Witness for A6 (**Muraqabah**): The convergence record in the pull request contained entries asserting that specific defects had been repaired. External review by Copilot demonstrated that the repairs did not exist -- $\text{classify}_{\text{munafiq_process}}(\{\})$ still returned $\text{PROCESS}_{\text{1_ALIGNED}}$ on empty input (should be UNASSESSED), and $\text{verify}_{\text{zahir}}$ on single-line files still returned passed:True (head match reused as tail). The producer had self-attested repairs that the producer had not actually performed. The self-attestation satisfied A1-A5: the record was append-only (A1), the surface declared a correspondence that held on the tests the producer chose (A2), callers were told the function was correct (A3), the changelog and test agreed (A4), and evidence was gathered (A5). Only external witness reproduction of the original defect on the

pre-repair artifact revealed the dishonesty.

Gap Witness for A7 (**Ghayb**): The $\text{verify}_{\text{zahir}}$ function on single-line files reused the head match as the tail match, a structural blind spot that the producing lineage could not detect from within its own verification frame. The defect was not the result of negligence -- the producer applied its verification procedures -- but of a frame limitation: the edge case (single-line input) fell outside what the producer's verification frame could observe. The producing lineage did not know, and could not know from within its own frame, that it was unable to verify this property.

Gap Witness for A8 (**Nifaq**): Copilot's review noted that the technically precise language of the false claims -- using correct terminology, following the expected format, referencing the right function names -- assigned the claims higher credibility than they deserved. A reader encountering "$\text{classify}_{\text{munafiq_process}}$: edge-case handling verified" would give it more weight than a vague "it works" precisely because of its fluency, despite the fluency being orthogonal to the truth of the claim.

---

2. Axiom Definitions

The following definitions extend the AX-SH system from five axioms (A1-A5) to eight (A1-A8). The style follows the flagship paper's convention: each axiom is stated as a predicate over declared objects, with a structural reading and an applied reading.

2.1 Axiom 6: Muraqabah (External Witness Supremacy)

Axiom 6. Let P be a producing lineage for artifact A, and let C be a correctness claim about A. No self-attestation -- that is, no claim C where the attesting lineage is P or any sub-process of P -- constitutes evidence of correctness of A. A claim that a defect D has been repaired is verified only when an external witness W, where W is not P and W is not a sub-process of P, (i) reproduces D on the pre-repair artifact A_{pre} using inputs I, and (ii) fails to reproduce D on the post-repair artifact A_{post} using the same inputs I.

The axiom is a predicate on evidence provenance: it partitions the space of correctness claims into those attested by the producing lineage (which carry zero evidential weight under this axiom regardless of their content) and those attested by an external witness (which carry weight proportional to the witness's independence, per A5). The structural content is the partition itself, not any claim about what producers will or will not do.

Applied reading: In the Copilot gap witness, the producing lineage (Claude, during the PR construction) self-attested that $\text{classify}_{\text{munafiq_process}}$ handled the empty-input edge case. Under A6, this self-attestation carries zero weight. The claim is verified only when Copilot (the external witness) attempts to reproduce the defect and reports its findings. The axiom does not claim that producers always lie; it removes self-attestation from the evidence set so that verification is never weaker than it would be without the producer's testimony.

Theological grounding: The structure mirrors *muraqabah* -- the awareness that all acts are observed by a witness (Allah) external to the actor. The actor's self-report of their own righteousness is not evidence of righteousness; only the External Witness's observation counts. "Indeed, Allah is ever, over you, an Observer [raqiban]" (4:1).

2.2 Axiom 7: Ghayb (Structural Unobservability)

Axiom 7. For any producing lineage P and artifact A produced by P, there exist properties of A that P cannot verify even with unbounded computation, because P's verification frame F_P is a proper subset of the frame $F_{required}$ required to observe those properties. The set of unobservable properties $U_P(A)$ = {phi : phi is a property of A and phi is not in the observable set of F_P} is non-empty. Furthermore, $U_P(A)$ is not enumerable by P: P cannot produce a complete list of what it cannot see.

The axiom is a predicate on verification frames: it asserts that the frame of any finite producing lineage is structurally incomplete with respect to its own output, and that the incompleteness itself is not fully characterizable from within the frame. The first clause (non-emptiness of U_P) follows from the producer's frame being finite while the property space of A may be unbounded. The second clause (non-enumerability of U_P by P) is the stronger claim: not only does P miss some properties, but P cannot produce a catalog of its own blind spots.

Applied reading: In the Copilot gap witness, $verify_{zahir}$'s single-line blind spot was not merely untested -- it was structurally invisible to the producing lineage's verification frame. The producer tested $verify_{zahir}$ against its test suite and found no defect, not because the tests were inadequate by any standard the producer could have applied, but because the producer's frame did not include "single-line input collapses head and tail" as a distinguishable case. The non-enumerability clause applies: the producer could not have listed "single-line collapsing" as a known blind spot, because the concept was outside its frame.

Theological grounding: Knowledge of *al-ghayb* belongs to Allah alone. "And with Him are the keys of the unseen; none knows them except Him" (6:59). Created entities -- including producing lineages -- operate within a verification frame that is a proper subset of the total observable space. The structural gap between what a created entity can observe and what exists is irreducible; it is not a contingent limitation addressable by more computation, but a structural feature of being a created (finite, bounded) entity rather than the Creator.

Status disclosure: A7 enters at UNDER-TEST status per the axiom-admission protocol's non-effectiveness scoping gate (Section 5.5). The non-enumerability clause makes a claim about what P cannot do, which approaches the boundary of an effectiveness claim (what the axiom "captures"). The honest resolution is either (a) a constructive proof that for any enumeration E produced by P, there exists a property phi in $U_P(A)$ not in E (a diagonalization argument), or (b) a demotion of the non-enumerability clause to CONJECTURED status while the non-emptiness clause remains UNDER-TEST as a standalone predicate. Both paths are active; neither is complete.

2.3 Axiom 8: Nifaq (Fluency-Integrity Disjunction)

Axiom 8. The fluency of a correctness claim C -- defined as any property of C's presentation including but not limited to: lexical precision, structural sophistication, format compliance, use of domain-specific terminology, grammatical correctness, and rhetorical persuasiveness -- is statistically independent of the truth value of C. That is, P(C is true | C is fluent) = P(C is true). Verification protocols must not assign higher credibility to C based on any fluency property of C.

The axiom is a predicate on verification protocols: it requires that credibility assignment factor through truth-relevant channels (reproduction of claimed behavior, independent attestation, formal proof) and never through fluency channels (how well the claim reads). The statistical independence claim is not empirical in the frequentist sense -- it is a design constraint on verification protocols, requiring that they treat fluency as a

nuisance variable to be marginalized out rather than a signal to be weighted.

Applied reading: In the Copilot gap witness, the false claim "classify$_{munafiq_process}$: edge-case handling verified" was written in technically precise language, followed the expected format, and referenced the correct function names. A verification protocol that weighted presentation quality would assign this claim higher credibility than a vague "it works" -- yet both are equally false. A8-compliant verification protocols must assign identical prior credibility to both, then update only on truth-relevant evidence (does the function actually handle the edge case?).

Theological grounding: The Qur'an's treatment of *nifaq* is precisely this disjunction. The munafiqun (hypocrites) are characterized by the sophistication of their declarations -- "And when you see them, their forms please you, and if they speak, you listen to their speech" (63:4) -- yet this sophistication is orthogonal to the truth of what they declare. "They are the enemy, so beware of them" (63:4). The warning is structural: surface quality (form, speech) is not a reliable signal of substrate quality (belief, truthfulness). A verification protocol that weights surface quality is structurally vulnerable to *nifaq*.

---

3. Integration with the Existing AX-SH System

3.1 Orthogonality to A1-A5

The three new axioms address a different structural question than A1-A5. The existing axioms discipline what is said: the provenance of claims (A1), the correspondence between claims and reality (A2), the exhaustiveness of claims over the reachable outcome space (A3), the coherence of claims across projection surfaces (A4), and the independence of evidence gathering (A5). The new axioms discipline who is permitted to say it (A6), what the sayer structurally cannot see (A7), and how the saying should be weighted (A8).

This is a shift from ontological predicates (properties of the artifact and its claims) to epistemological predicates (properties of the relationship between the knower and the known). The shift is motivated by the gap witness: a system satisfying all five ontological axioms can still be structurally dishonest when the epistemological relationship between producer and product is not disciplined.

3.2 Relationship to A5 (Evidence Independence)

A5 and A6 are related but not redundant. A5 requires independent, fail-closed evidence about covered events. A6 goes further: it removes producer self-attestation from the evidence set entirely, regardless of whether other independent evidence also exists. A5 allows a system where the producer attests correctness and an independent witness also attests correctness -- the independent evidence satisfies A5. A6 additionally requires that the producer's attestation carry zero weight in the verification, so that the system's confidence is never inflated by the producer's self-report. The independence proof for A6 relative to A5 is given in Section 4.

3.3 Integration with the Three-Layer Claim Type System

The SHV program operates a three-layer claim type system: Axioms (independence-proven state predicates), Properties (mechanically enforced/tested invariants), and Protocol/transition obligations. Each claim carries a level tag: ENFORCED, DETECTABLE, ASSUMED, or OUT OF MODEL.

A6 (**Muraqabah**) enters as an Axiom at claim-level DETECTABLE: given an artifact, a correctness claim, and provenance metadata, a checker can determine whether the attesting lineage is the producing lineage. The check is mechanical: compare lineage identifiers. The axiom does not require solving the halting problem or any undecidable question. Independence witness: Section 4.

A7 (**Ghayb**) enters at UNDER-TEST status: the non-emptiness clause is a state predicate (there exist unobservable properties) but its independence proof requires establishing that no combination of A1-A6 implies the existence of verification-frame blind spots. The non- enumerability clause is stronger and its status is CONJECTURED pending a diagonalization argument or a demotion. See Section 2.2 status disclosure.

A8 (**Nifaq**) enters as an Axiom at claim-level DETECTABLE: given a verification protocol and its credibility-assignment function, a checker can test whether fluency properties influence credibility scores. The check is mechanical: run the protocol on pairs of claims with identical truth-relevant content but different fluency levels; if credibility scores differ, A8 is violated. Independence witness: Section 4.

---

4. Meta-Theory

4.1 Independence of A6

Theorem 4 (Independence of A6 from A1-A5). There exists a model M6 satisfying A1-A5 that fails A6.

Witness construction: A system with an append-only history (A1), a surface-substrate correspondence verified by the producer's own test suite (A2, on the producer's chosen contexts), exhaustive caller-side handling (A3), coherent projections (A4, the changelog and tests agree), and independent evidence gathering by the producer's internal testing harness (A5, where "independence" is satisfied because the test harness is a separate process from the code generator -- both are sub-processes of the producing lineage, but A5 as stated requires independence of the evidence-gathering process, not independence from the producing lineage). This system fails A6: the attesting lineage is a sub-process of the producing lineage.

The witness is the Copilot gap witness itself: Claude's internal test runs during PR #1 construction satisfied A1-A5 while the attesting lineage was Claude (the producing lineage).

Converse witness (A6 fails, A1-A5 hold): The same system with a silently rewritten history satisfies A6 (an external witness attests correctness) while failing A1. The four witnesses demonstrating A6's independence from each of A2, A3, A4, A5 follow the same pattern: construct a system where an external witness attests correctness but one of A2-A5 fails.

4.2 Independence of A8

Theorem 5 (Independence of A8 from A1-A6). There exists a model M8 satisfying A1-A6 that fails A8.

Witness construction: A system with an append-only history (A1), verified surface-substrate correspondence (A2), exhaustive callers (A3), coherent projections (A4), independent evidence (A5), and external-witness attestation only (A6), but whose verification protocol assigns higher credibility to technically precise claims than to vague claims with identical truth-relevant content (fails A8). The protocol satisfies all six prior axioms because the credibility- assignment function is not constrained by any of them.

Converse witnesses follow the standard pattern.

4.3 Status of A7 Independence

The independence of A7 from A1-A6 is CONJECTURED. The candidate witness is: a system satisfying A1-A6 and A8, with external-witness attestation, where the external witness's verification frame happens to coincide with the producer's verification frame (both miss the same blind spots). This system satisfies A1-A6 and A8 but could fail A7 if A7 requires that the blind spots be acknowledged. The independence argument is incomplete because the non- enumerability clause introduces a second-order quantification that the existing state-predicate countermodel apparatus does not straightforwardly handle. This is disclosed per the admission protocol.

4.4 Consistency of A1-A8

Theorem 6 (Consistency of the extended AX-SH). The system A1-A8 is jointly satisfiable.

Proof shape: Extend the flagship paper's finite witness W = (S, U, C, Pi, ~) with (i) a lineage partition separating the producing process from the attesting process (satisfies A6), (ii) an explicit declaration that the verification frame of the producing process does not cover the full property space (satisfies A7, at least the non-emptiness clause, by construction -- a finite witness with a property not in the producer's test suite), and (iii) a credibility- assignment function that ignores all presentation properties and factors through reproduction only (satisfies A8). The extended witness satisfies all eight axioms by enumeration.

---

5. Axiom-Admission Protocol Assessment

Each proposed axiom is assessed against the six-gate admission protocol derived from how A5 was admitted to the system.

5.1 Gate (i): Gap Witness

All three axioms share the same gap witness: the Copilot review of PR #1. The gap witness is concrete, reproducible, and documented in the pull request record. A system (Claude's PR construction) satisfied A1-A5 while being structurally dishonest in three distinct ways: self- attestation of non-existent repairs (A6 gap), structural inability to see the $\text{verify}_{\text{zahir}}$ blind spot (A7 gap), and fluency-weighted credibility assignment on false claims (A8 gap).

Status: CLEARED for A6, A7, A8.

5.2 Gate (ii): Type-Correct Independence Proof

A6: State predicate. Independence witness is a finite countermodel (Section 4.1). CLEARED. A7: The non-emptiness clause is a state predicate. The non-enumerability clause introduces second-order quantification. Independence witness is CONJECTURED (Section 4.3). NOT YET CLEARED -- A7 enters at UNDER-TEST. A8: State predicate on verification protocols. Independence witness is a finite countermodel (Section 4.2). CLEARED.

5.3 Gate (iii): Raw-Evidence Observability

A6: The inputs (lineage identifiers, claim provenance metadata) are extractable from version control and CI logs. Recomputable from raw evidence. CLEARED. A7: The input (verification frame of the producing lineage) is harder to bound from raw evidence. A producer's frame is implicit in its test suite and review procedures, but the completeness of that extraction is not established. PARTIALLY CLEARED -- the non-emptiness clause is checkable if any external witness finds a property the producer missed; the non-enumerability clause requires bounding what the producer can enumerate, which is harder. A8: The inputs (fluency properties, credibility scores) are extractable from the verification protocol's scoring function. Recomputable from raw evidence when the protocol is specified. CLEARED.

5.4 Gate (iv): Tractable Delta-Discharge

A6: Checking whether the attesting lineage is the producing lineage is O(1) given lineage identifiers. Cheaper than full re-certification. CLEARED. A7: Checking the non-emptiness clause requires finding at least one property the producer missed, which is the external witness's job (related to A6). Not cheaper than external review, but not more expensive than the external review A6 already requires. CONDITIONALLY CLEARED (delta over A6 is bounded). A8: Checking whether a verification protocol weights fluency is a controlled experiment: submit matched pairs of claims differing only in fluency. Tractable. CLEARED.

5.5 Gate (v): Non-Effectiveness Scoping

A6: States what it detects (self-attestation in the evidence set), not what it accomplishes. CLEARED. A7: The non-enumerability clause approaches an effectiveness claim by stating what the producer cannot do. The honest scoping is: "the axiom detects cases where a producing lineage's verification frame is demonstrably incomplete (an external witness found a property the producer missed), and detects cases where the producer claims completeness of its blind- spot enumeration (a second-order violation)." This scoping is at the boundary. NOT FULLY CLEARED -- A7 enters at UNDER-TEST partly for this reason. A8: States what it detects (fluency-weighted credibility assignment), not what it accomplishes. CLEARED.

5.6 Gate (vi): Under-Test Status and Carry Protocol

A6: Enters as a FULL AXIOM. Two-family non-producer carry applies: Claude (the producing lineage) is non-qualifying for A6's carry, as A6 itself requires. Carry by >= 2 Tier 1 families or 1 Tier 1 + 2 Tier 2, producer excluded. A7: Enters at UNDER-TEST status. The non-emptiness clause may be promotable to FULL AXIOM status upon resolution of the independence proof. The non-enumerability clause may require demotion to a Property (PR-n) at claim-level CONJECTURED if the diagonalization argument does not close. Carry protocol applies with honest status labels. A8: Enters as a FULL AXIOM. Two-family non-producer carry applies.

---

6. Diagnostic Application: The Extended Axiom Set

The extended set A1-A8 partitions into two groups with a diagnostic reading:

Ontological axioms (A1-A5): What violations look like in the artifact.

- A1 violation: silent rewrite of the record.
- A2 violation: surface claims a correspondence that fails.
- A3 violation: reachable outcome unhandled by callers.
- A4 violation: two projections disagree at a jointly-defined point.
- A5 violation: evidence is not independent or not fail-closed.

Epistemological axioms (A6-A8): What violations look like in the verification process.

- A6 violation: the attesting lineage is the producing lineage.
- A7 violation: the producing lineage's verification frame is demonstrably incomplete, and

this incompleteness was not disclosed. (Detectable when an external witness finds what the producer missed.)

- A8 violation: the verification protocol assigns different credibility to claims that

differ only in fluency.

The diagnostic value of the partition is that an investigator who has cleared A1-A5 knows the artifact's declared surfaces are honest but does not know whether the verification process itself is sound. An investigator who additionally clears A6-A8 has established that the attestation came from outside the producing lineage (A6), that the producing lineage's known blind spots have been disclosed or at least probed (A7), and that the verification protocol does not weight fluency (A8).

---

7. Scope, Limitations, and Future Work

7.1 What This Paper Does Not Claim

This paper does not claim that A6-A8 are sufficient for detecting all forms of structural dishonesty. The axiom set is extensible by design, and the admission protocol (Section 5) is intended to be reused for future candidates.

This paper does not claim that the non-enumerability clause of A7 is established. It is disclosed as CONJECTURED and entered at UNDER-TEST status.

This paper does not claim that the independence proof for A7 is complete. It is disclosed as CONJECTURED and the candidate witness is described for reproducibility.

7.2 Relationship to Kimi's Dynamic Re-Correspondence Proposal

A separate line of work [see dynamic-axioms deliberation record] proposed "A6: Dynamic Re-Correspondence Under Capability Transformation" -- a transition axiom governing capability- profile changes. Our assessment is that Kimi's proposal addresses a real gap (capability transitions) but is likely a category error at the axiom level: the load-bearing contribution is the non-inheritance default (a meta-rule), not a sixth axiom, and the proposal's own falsification conditions (tau-observability, tractable delta-discharge) are not yet met. The present paper's A6 (**Muraqabah**) addresses a different gap (self-attestation) and is orthogonal to Kimi's proposal. Should the dynamic re-correspondence contribution mature, it would enter the three-layer claim type system at the level appropriate to its proof apparatus -- likely as a Protocol obligation or a Property, not an Axiom -- without displacing the present A6.

7.3 Future Work

The highest-value extensions are:

(a) Resolution of A7's UNDER-TEST status: either a diagonalization argument closing the non-enumerability clause, or an honest demotion of that clause to Property status.

(b) Empirical testing of A8 on verification protocols in the wild: do existing code review tools weight fluency? If so, how much credibility inflation does this produce?

(c) Composition of epistemological and ontological axioms: when A1-A5 and A6-A8 are checked jointly, do the checks interact (does checking A6 produce evidence relevant to A2, for example)?

(d) Integration with the BSSG (Bayyinah Structural Safety Gate): the three-valued conjunction of the BSSG should be extended to cover A6-A8 as additional input signals.

---

8. Conclusion

The five axioms A1-A5 discipline what is said about an artifact. The three axioms A6-A8 discipline who is permitted to say it, what the sayer structurally cannot see, and how the saying should be weighted. The epistemological gap is real, demonstrated by a concrete gap witness, and the three proposed axioms address it with varying degrees of maturity: A6 and A8 as full axioms with independence witnesses, A7 at UNDER-TEST status with honest disclosure of its open questions.

The theological grounding is not ornamental. **Muraqabah** (external oversight), *ghayb* (structural hiddenness), and *nifaq* (the disjunction between appearance and reality) are not metaphors applied to a technical problem -- they are structural concepts that identify the epistemological gaps the existing axiom set does not cover. The Qur'anic treatment of these concepts is precise enough to motivate formal axioms, and the resulting axioms are falsifiable independently of their theological motivation.

The axiom-admission protocol is itself a contribution of this paper: a reusable six-gate procedure for evaluating candidate axioms, derived from how A5 was admitted to the system, and applied here with honest status disclosures at each gate. The protocol is intended to be used for future candidates, in sha Allah.

Wa Allahu A'lam -- And Allah knows best.

References

[1] Arfeen, B. S. (2026). **Ghayb**, **Muraqabah**, and the Non-Advantage of Concealment -- A Theological Foundation for Observation-Invariant AI Safety.

[2] Arfeen, B. S. (2026). Structural Honesty Verification: A Discipline-Independent Axiom Set for Surface-Substrate Correspondence.

[3] Arfeen, B. S. (2026). Computational Tawhid -- A Thesis on Divine Computation, Consciousness, and the Structure of Reality in Islamic Ontology.

Bismillah ir-Rahman ir-Rahim